



GOLD SPONSORS

TECHNOLOGY
INNOVATION
DATA
KNOWLEDGE



SILVER SPONSORS



BRONZE SPONSOR



STRATEGIC PARTNER

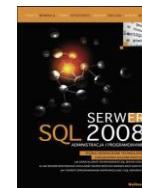
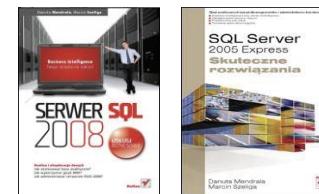


Image classification with Microsoft Cognitive Toolkit

marcin.szeliga@datacommunity.pl

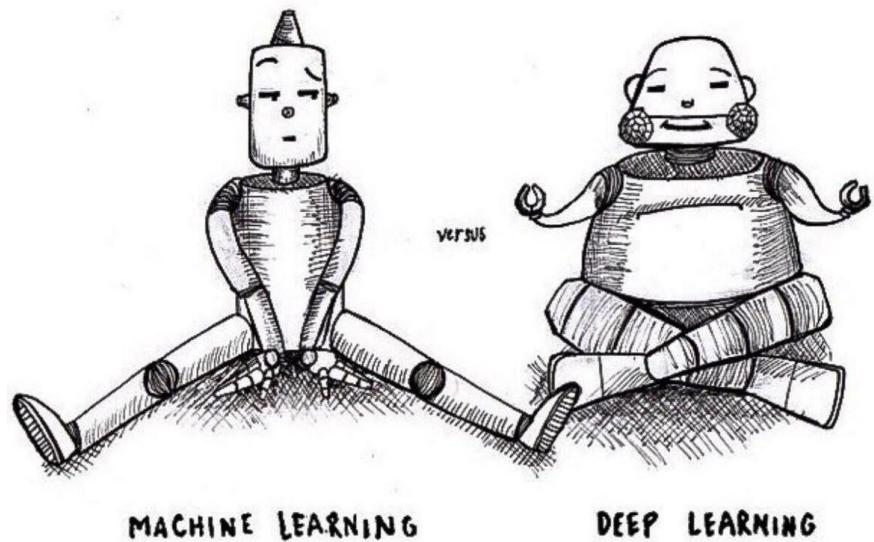
Marcin Szeliga

- Data philosopher
- 20 years of experience with SQL Server
- Data Platform MVP & MCT
- Microsoft Certified Solutions Expert
 - Data Management and Analytics
 - Cloud Platform and Infrastructure
 - Business Intelligence
- Microsoft Certified Solutions Developer
 - Azure Solution Architect

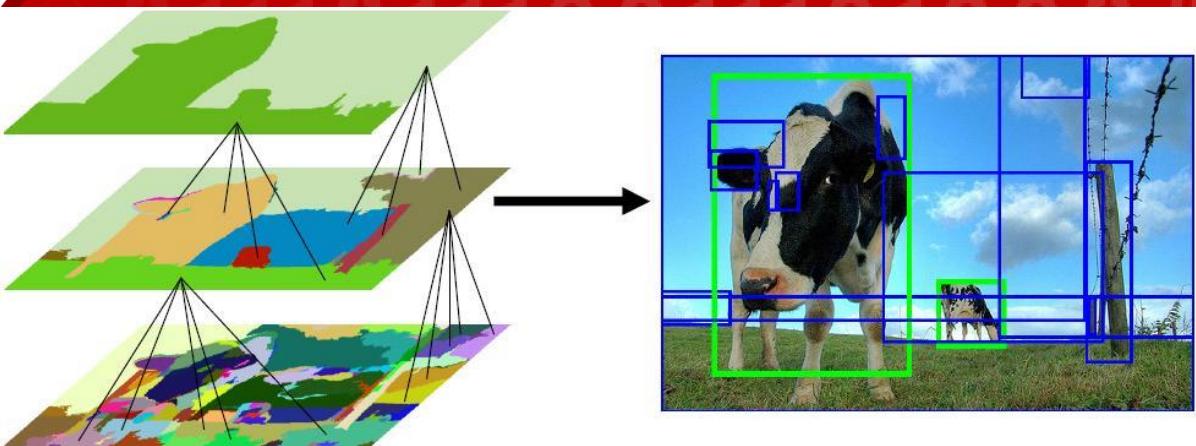


Agenda

- Why image classification is so difficult?
- Nearest Neighbor Classifier
- Linear classification
- Regular Neural Networks
- Convolutional Neural Networks

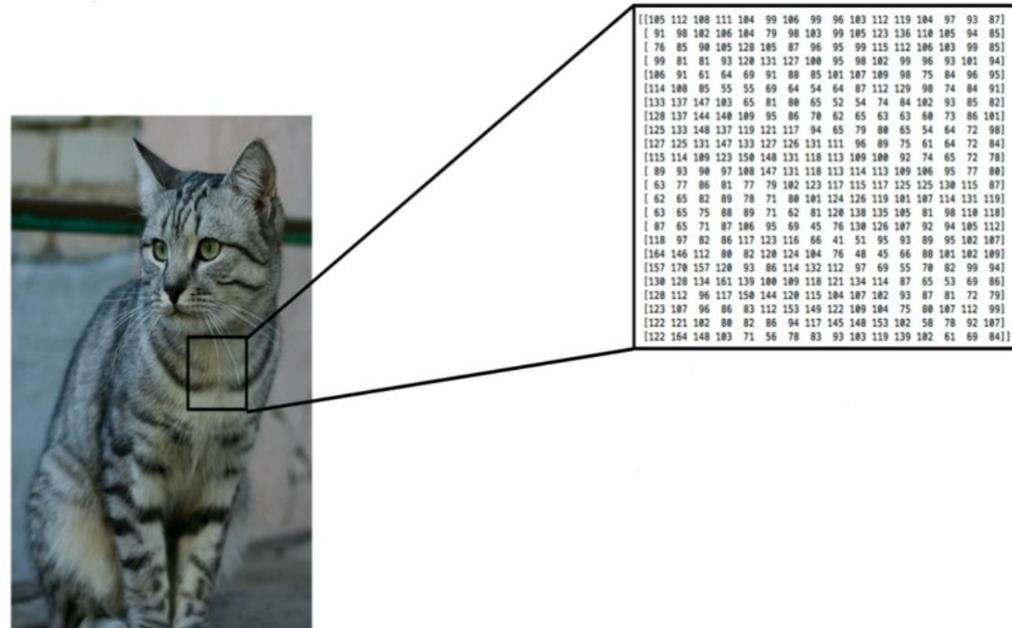


Why image classification is so difficult?



To a computer an image is represented as one large 3-dimensional array of numbers

- Cat image is 248 pixels wide, 400 pixels tall, and has three color channels Red,Green,Blue
- The image consists of $248 \times 400 \times 3$ numbers = 297,600
- Each number is an integer that ranges from 0 (black) to 255 (white)



Challenges

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation



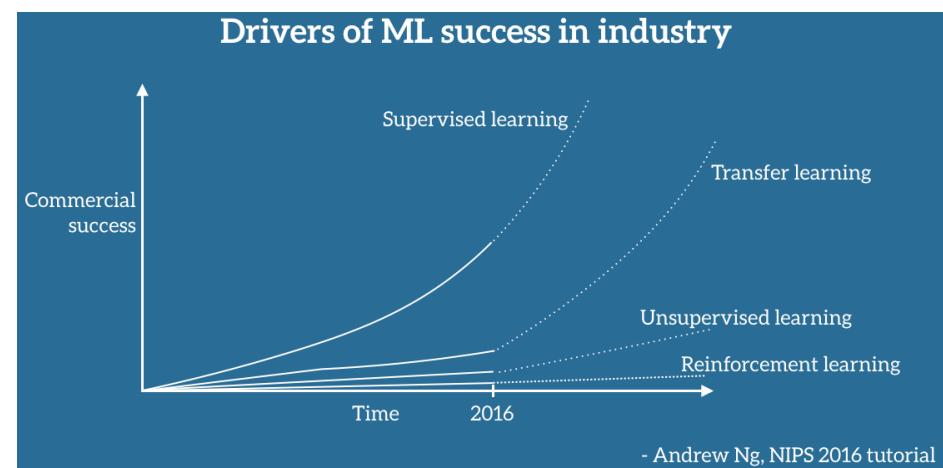
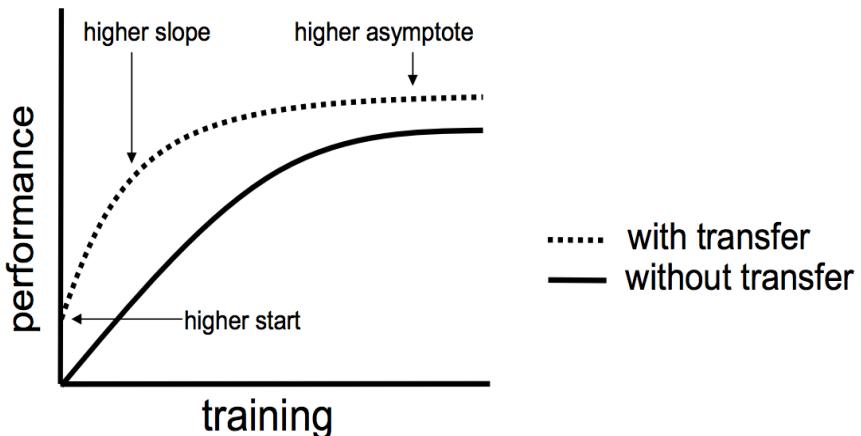
Good image classification model must be invariant to the cross product of all these variations

Fast solution

- Custom vision: <https://www.customvision.ai/>
- Easily customize your vision models that fit perfectly with your unique use case
- Bring your own labeled images, or use Custom Vision to quickly add tags to any unlabeled images
- Use your labeled images to teach Custom Vision the concepts you care about
- Use simple REST API calls to quickly tag images with your new custom computer vision model

Quick introduction to transfer learning

- Machine learning method where a model developed for a task is reused as the starting point for a model on a second task
- Very popular in deep learning given the enormous resources required to train deep learning models
- Use a deep learning model pre-trained for a large and challenging image classification task such as the ImageNet 1000



A screenshot of a web browser window titled "Custom Vision". The address bar shows the URL "https://customvision.ai/projects". The main content area is titled "Projects" and displays a message "You have no projects yet, create one!". Below this message is a large blue button labeled "New Project" with a plus sign icon. The browser interface includes standard navigation buttons (back, forward, search) and a user profile icon.

Projects



New Project

You have no projects yet, create one!



Microsoft Cognitive Toolkit

- A free, easy-to-use, open-source, commercial-grade toolkit that trains deep learning algorithms to learn like the human brain
 - <https://www.microsoft.com/en-us/cognitive-toolkit/>
- Highly optimized, built-in components
 - Automatic hyperparameter tuning
 - Built-in readers optimized for massive datasets
- Efficient resource usage
 - Parallelism with accuracy on multiple GPUs/machines via 1-bit SGD and Block Momentum

Filter

The Microsoft Cognitive Toolkit

07/31/2017 • 2 minutes to read • Contributors  all

[Feedback](#) [Edit](#) [Share](#)

Theme [Light](#)

Overview

- What's new
- > Getting Started
- > Tutorials & Examples
- > Train/Develop
- > Evaluate/Deploy
- > API Reference
- > How To
- > CNTK Source Code & Development
- > BrainScript
- > Resources
- FAQ
- Feedback

The Microsoft Cognitive Toolkit (CNTK) is an open-source toolkit for commercial-grade distributed deep learning. It describes neural networks as a series of computational steps via a directed graph. CNTK allows the user to easily realize and combine popular model types such as feed-forward DNNs, convolutional neural networks (CNNs) and recurrent neural networks (RNNs/LSTMs). CNTK implements stochastic gradient descent (SGD, error backpropagation) learning with automatic differentiation and parallelization across multiple GPUs and servers.

This video provides a high-level overview of the toolkit. In addition, Microsoft offers an introductory course to deep learning with CNTK, [Deep Learning Explained](#).

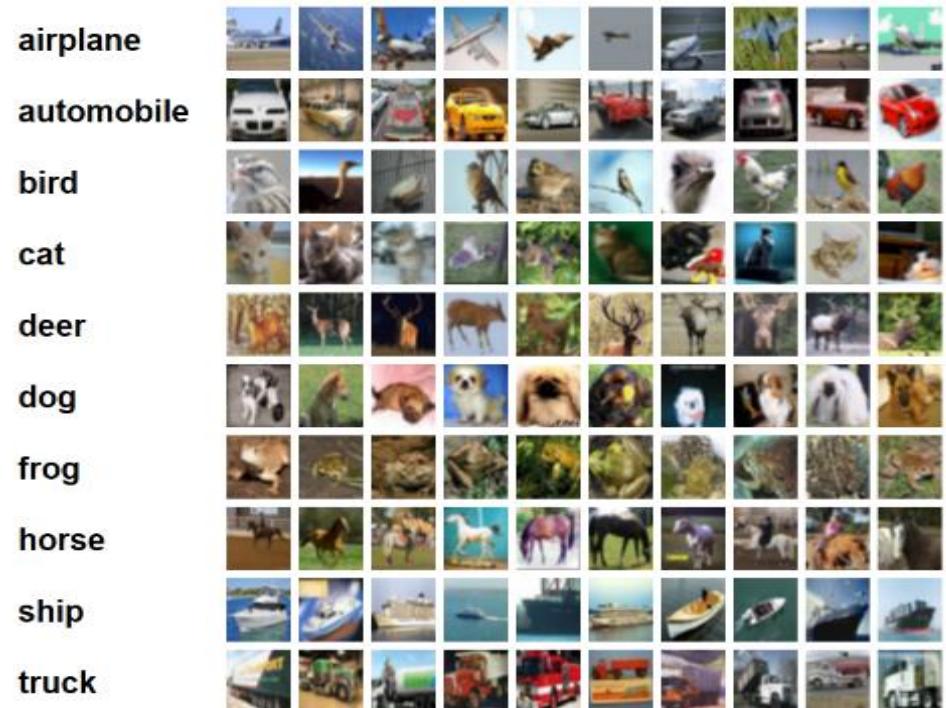
The latest release of CNTK is [2.3](#).

CNTK can be included as a library in your Python, C#, or C++ programs, or used as a standalone machine-learning tool through its own model description language (BrainScript). In addition you can use the CNTK model evaluation functionality from your Java programs.

CNTK supports 64-bit Linux or 64-bit Windows operating systems. To install you can either choose pre-compiled binary packages, or compile the toolkit from the source provided in [GitHub](#).

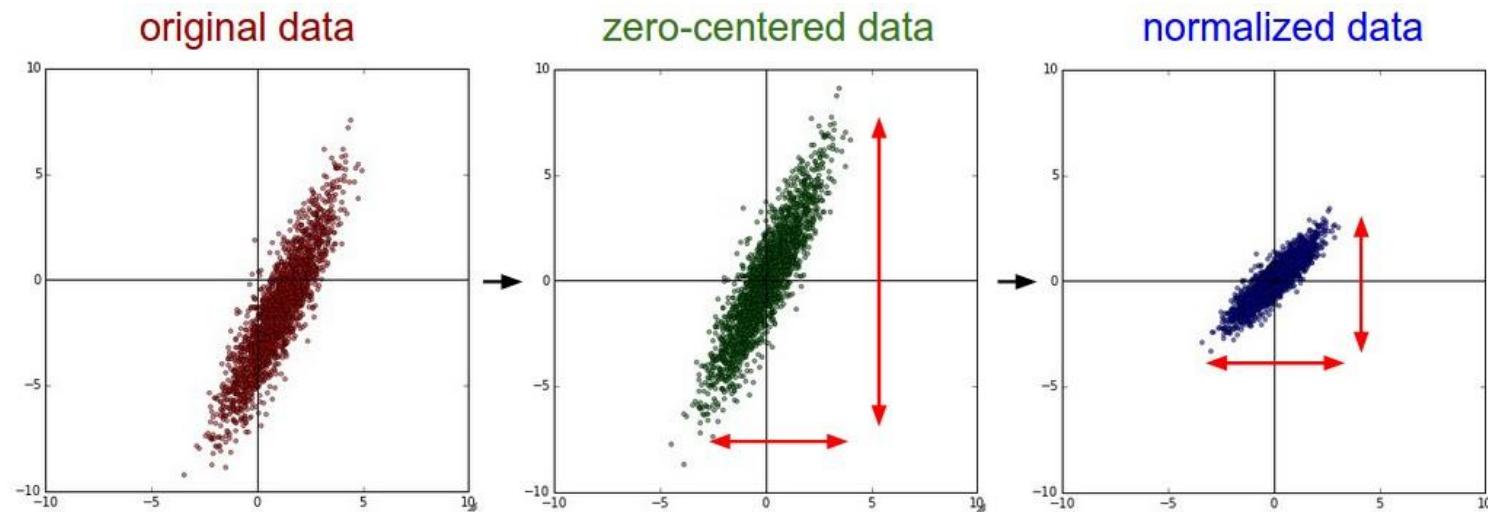
CIFAR-10 dataset

- Consists of 60 000 32x32 colour images in 10 classes, with 6 000 images per class
- Classes are completely mutually exclusive
- 50 000 training images and 10 000 test images
- <https://www.cs.toronto.edu/~kriz/cifar.html>



Data Preprocessing

- Mean subtraction is the most common form of preprocessing
 - Subtract a single value from all pixels or do so separately across the three color channel
- Normalization scale values to a specific range



```
C:\local\CNTK-2-2\cntk\Scripts>cntkpy35.bat
```

```
(C:\local\Anaconda3-4.1.1-Windows-x86_64\envs\cntk-py35) C:\local\cntk-2-2\cntk\Scripts>jupyter notebook
[I 07:33:16.928 NotebookApp] Serving notebooks from local directory: C:\local\cntk-2-2\cntk\Scripts
[I 07:33:16.929 NotebookApp] 0 active kernels
[I 07:33:16.929 NotebookApp] The Jupyter Notebook is running at: http://localhost:8888/?token=815c7c37f92035a3a9afe254c6d8c01e8c469680ebfa9cf4
[I 07:33:16.929 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 07:33:16.934 NotebookApp]
```

Copy/paste this URL into your browser when you connect for the first time,

to login with a token:

```
http://localhost:8888/?token=815c7c37f92035a3a9afe254c6d8c01e8c469680ebfa9cf4
```

```
[I 07:33:17.710 NotebookApp] Accepting one-time-token-authenticated connection from ::1
```

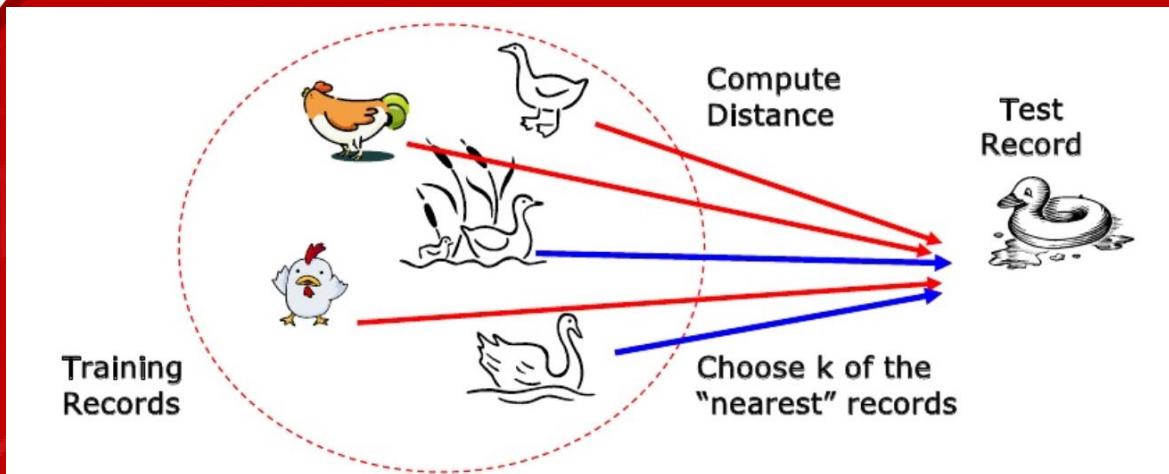
```
[I 07:33:26.815 NotebookApp] Kernel started: 93368b60-504d-4635-89e2-e79c2e8d75a7
```

```
[I 07:33:28.066 NotebookApp] Adapting to protocol v5.1 for kernel 93368b60-504d-4635-89e2-e79c2e8d75a7
```

```
[I 07:33:42.914 NotebookApp] Saving file at /LogisticRegression.ipynb
```

```
[I 07:33:48.087 NotebookApp] Kernel started: 6e0d6ed2-fb9c-49d9-a55c-ad27
```

Nearest Neighbor Classifier



How to compare two images?

- L1 distance - compare the images pixel by pixel and add up all the differences

Train image		Test image		Pixel-wise absolute difference	
30	95	50	200	20	105
60	43	33	140	27	97

$$- L1 = d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p| = 249$$

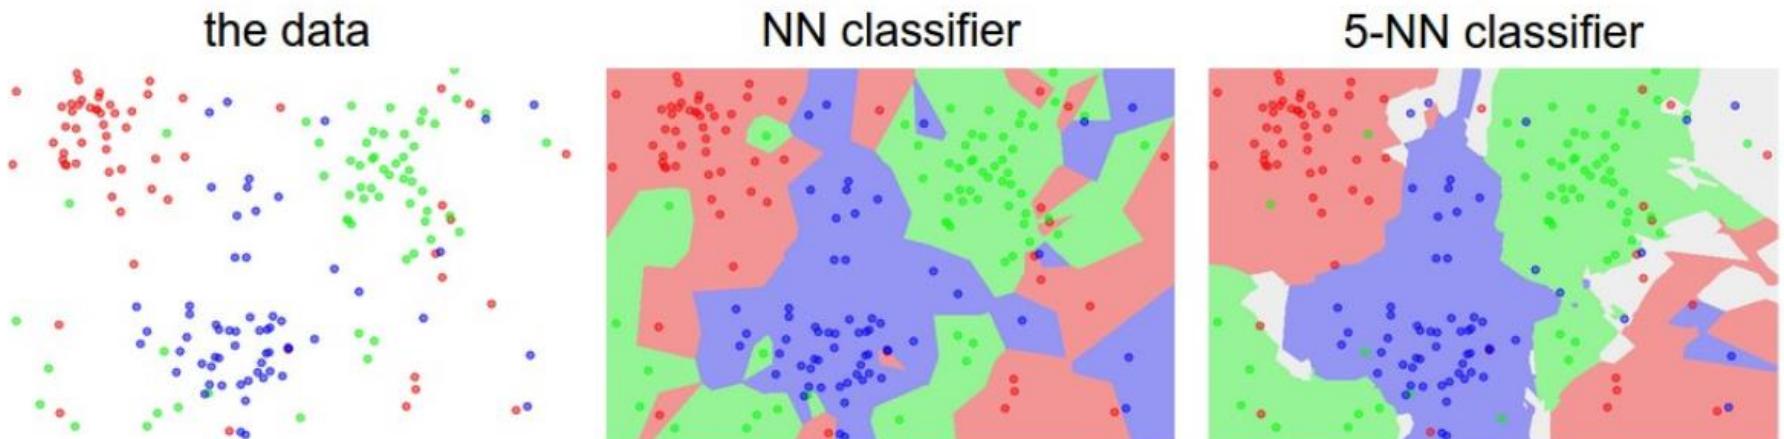
- L2 distance has the geometric interpretation of computing the Euclidean distance between two vectors

$$- L2 = d_1(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2} = 147$$

- Prefers many medium disagreements to one big one

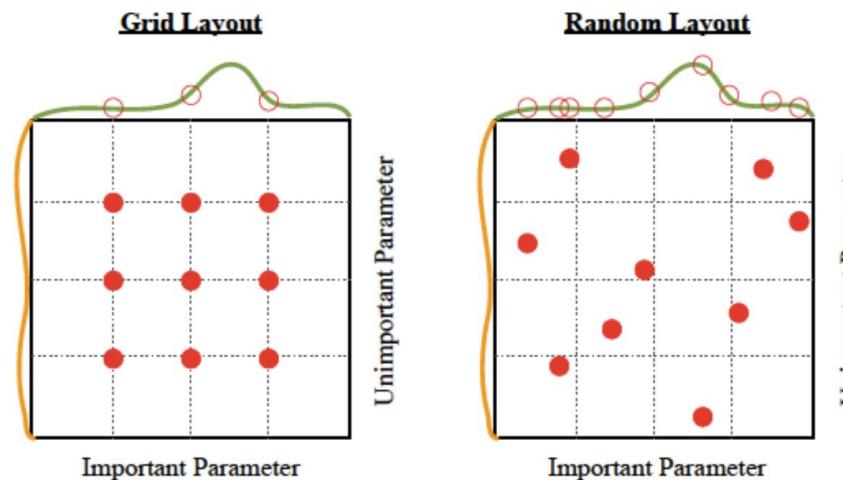
K-Nearest Neighbor Classifier

- Image can be labeled as the more similar train image
- Nonparametric models can capture more subtle aspects of the data, but are huge and slow
- What value of k should you use?

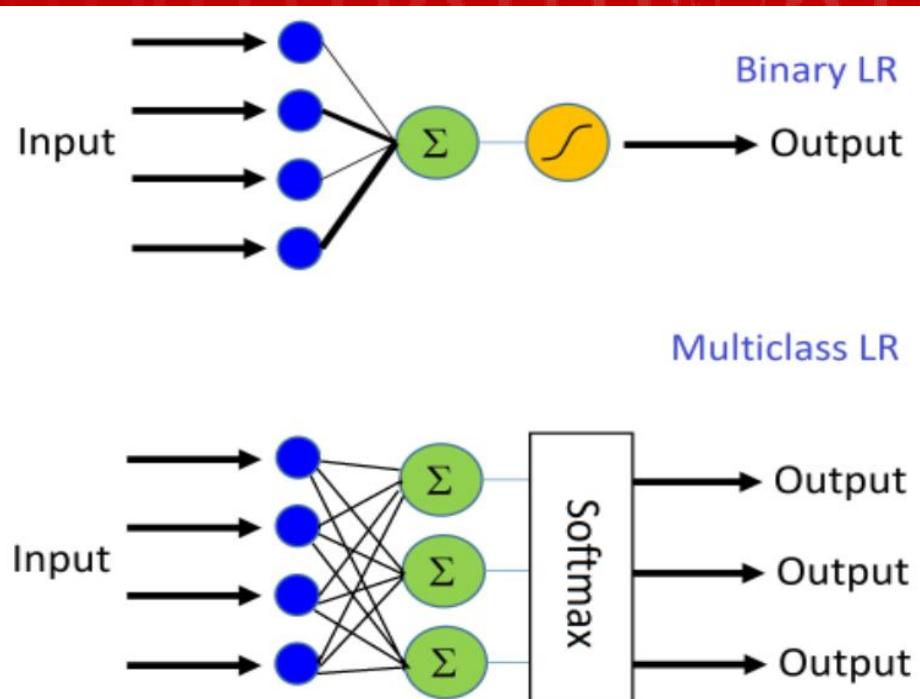


Hyperparameters

- Distance function and number of nearest neighbors are model hyperparameters
- We should try out many different values and see what works best
- But we cannot use the train set neither test set for the purpose of tweaking hyperparameters
- The idea is to split training set in two:
 - Slightly smaller training set
 - And a validation set

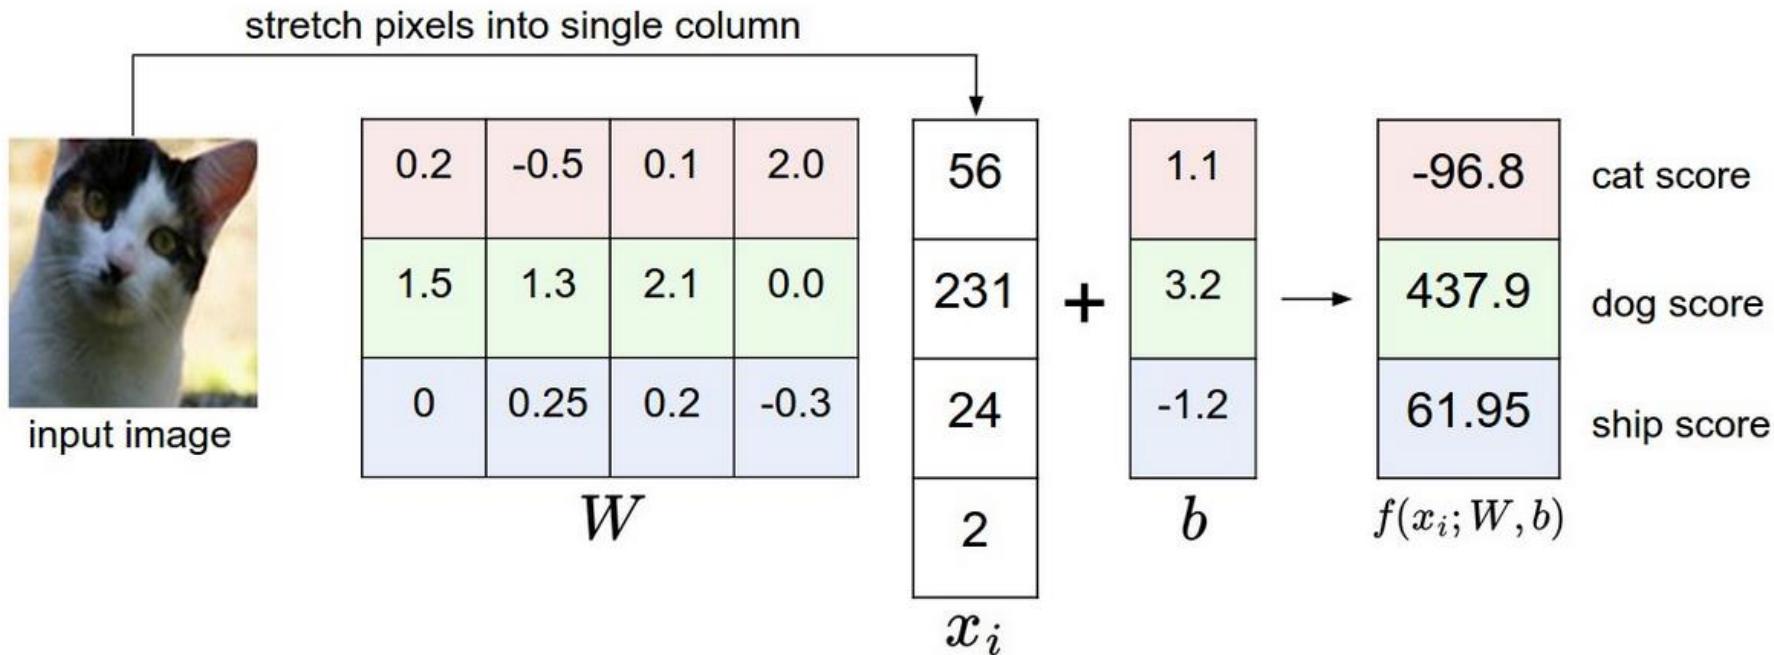


Linear Classification



Parameterized models

- We need a score function that maps the pixel values of an image to confidence scores for each class
- The simplest is linear regression: $y = f(x) = Wx_1 + b$



Template matching

- Each row of W corresponds to a template
- The score of each class for an image is obtained by comparing each template with the image using dot product one by one to find the one that fits best
- Like Nearest Neighbor, but instead of having thousands of training images we are only using a single image per class

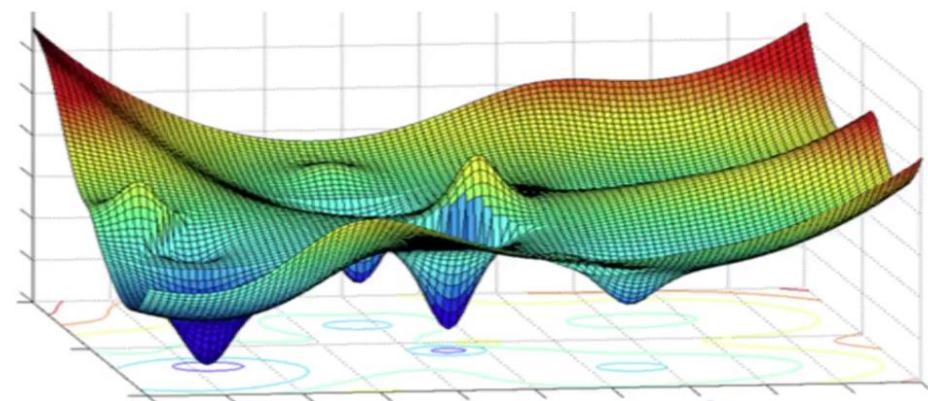


Loss function

- The task is to find a function that minimizes the difference between predicted and actual values
- The loss function quantifies our unhappiness with predictions on the training set

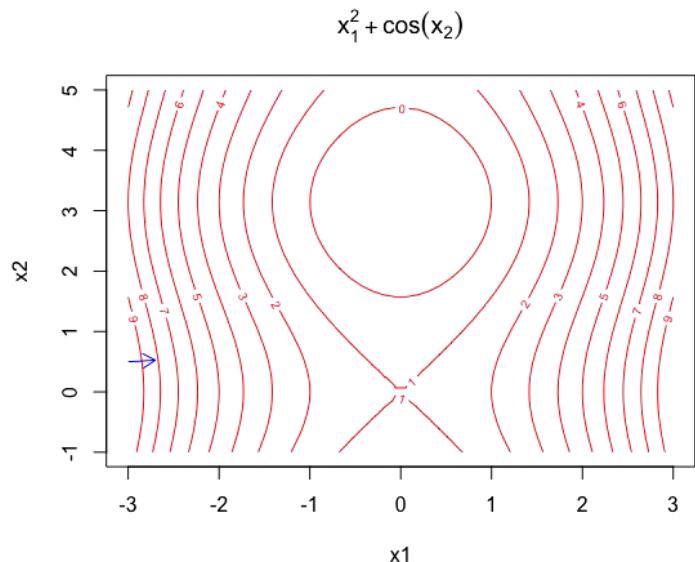
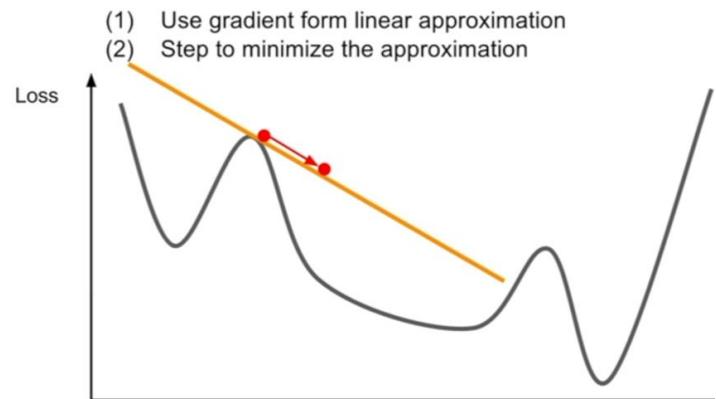
$$se = \sum_{j=0}^9 (y_j - p_j)^2$$

$$ce = - \sum_{j=0}^9 y_j \log(p_j)$$



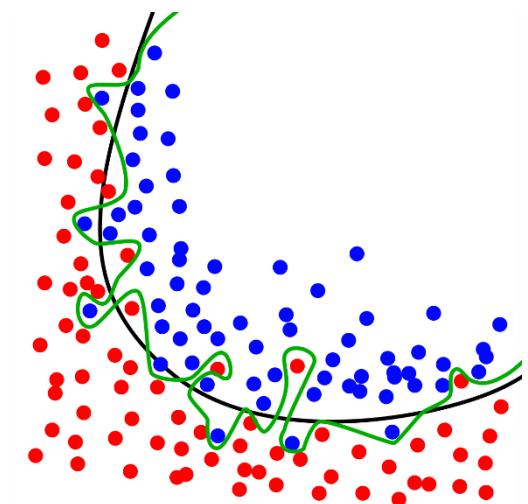
Stochastic gradient descent

- Loss function lets us quantify the quality of any particular set of weights W
- Goal of optimization is to find W that minimizes the loss function
- We can compute the best direction along which we should change our weight vector that is mathematically guaranteed to be the direction of the steepest descend



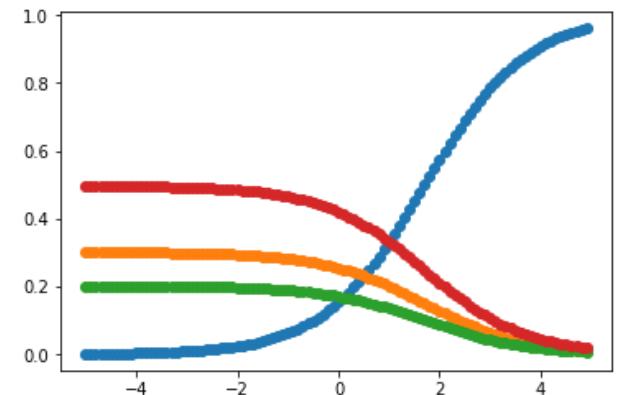
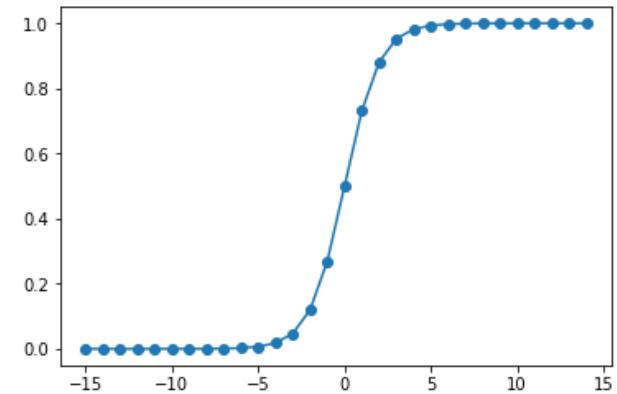
Overfitting

- For a given formula loss function can be very small for train set
- To prevent overfitting the model should be generalized
 - We can add regularization term to loss function
 - L2 regularization (Ridge) adds squared magnitude of weights as penalty term to the loss function
 - L1 regularization (Lasso) adds absolute value of magnitude of weights as penalty term to the loss function
 - We also can
 - Stop training early
 - Dropout random parameters per minibatch
 - But adding more images yields the best results
 - Use data augmentation



Probability distribution

- Our goal is try to predict the probability that a given example belongs to the specific class
- We need a s-shaped function that squashes values into the range $[0,1]$ so that we may interpret them as a probability
- Model has to search for a value of loss function so that the probability is large when x belongs to the right class and small when x belongs to any other classes



Sigmoid and Softmax explained

class	output	exp	sigmoid	softmax	actual
cat	-4.64	0.01	0.009535211	0.00	0
dog	-2.68	0.07	0.064288227	0.00	0
plane	2.87	17.55	0.946089038	0.98	1
fish	-2.08	0.13	0.11129408	0.01	0
building	-2.26	0.10	0.094315593	0.01	0
sum	-8.80	17.86	1.23	1.00	

$=\text{EXP}(\text{output})$

$=\exp/(1+\exp)$

$=\exp/\sum(\exp)$

One-hot-encoded

Home  x localhost:8888/notebooks/Image classification with Microsoft Cognitive Toolkit (autosaved) 140% ... Szukaj Logout

Jupyter Image classification with Microsoft Cognitive Toolkit (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

Code

```
Data directory is ..\Examples\Image\DataSet\cifar-10
```

In [3]: # Create the train and test readers
reader_train = create_reader(os.path.join(data_path, 'train_map.txt'),
 os.path.join(data_path, 'CIFAR-10_mean.xml'), T
reader_test = create_reader(os.path.join(data_path, 'test_map.txt'),
 os.path.join(data_path, 'CIFAR-10_mean.xml'), E
Reading map file: ..\Examples\Image\DataSet\cifar-10\train_map.txt
Reading mean file: ..\Examples\Image\DataSet\cifar-10\CIFAR-10_mean.xml
Reading map file: ..\Examples\Image\DataSet\cifar-10\test_map.txt
Reading mean file: ..\Examples\Image\DataSet\cifar-10\CIFAR-10_mean.xml

In []: #
Train and evaluate the network.

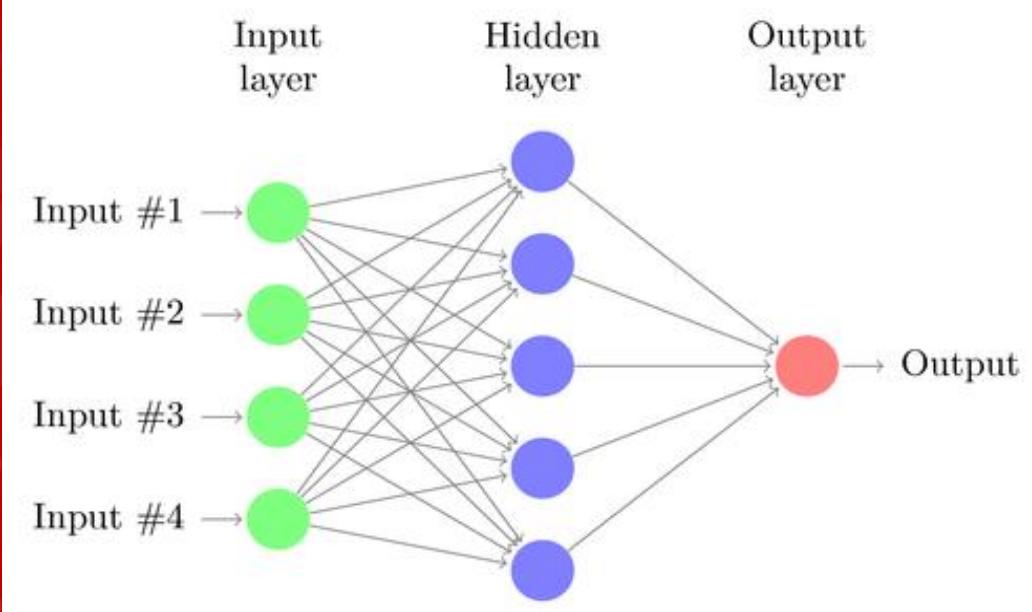
def train_and_evaluate(reader_train, reader_test, max_epochs, model_func):
 # Input variables denoting the features and label data
 input_var = C.input_variable((num_channels, image_height, image_width))
 label_var = C.input_variable((num_classes))

 # Normalize the input
 feature_scale = 1.0 / 256.0
 input_var_norm = C.element_times(feature_scale, input_var)

 # apply model to input
 z = model_func(input_var_norm, out_dims=10)

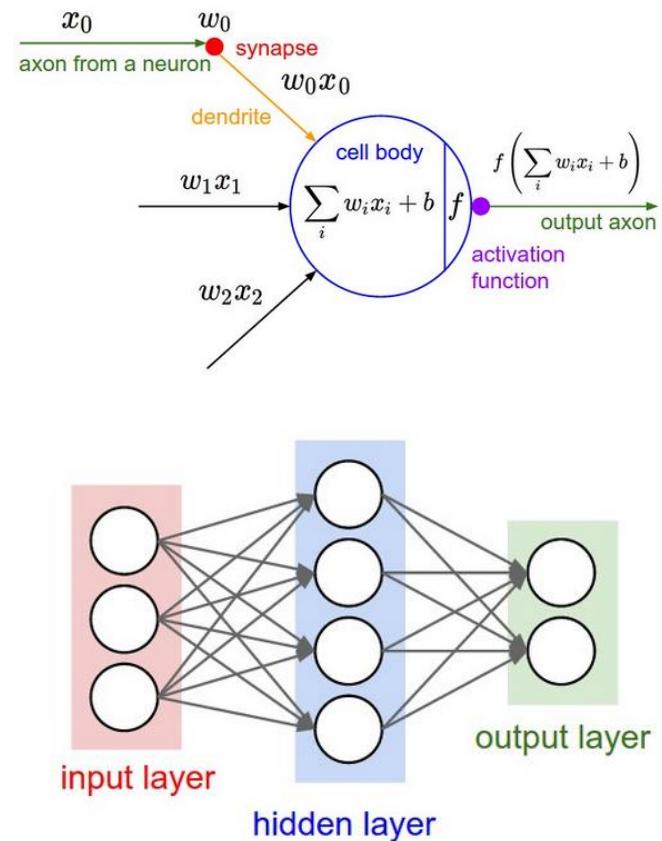
 #
 # Training action
 #

Regular Neural Networks



Artificial neural networks

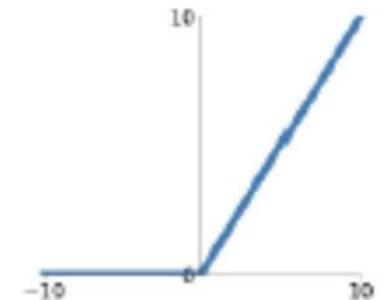
- Basic computational unit in is a neuron (perceptron)
- Input signals interact multiplicatively with outputs of the other neurons based on weight
- Weight signals get summed and bias is added
- Activation function applies the non-linearity
- Think of :
 - Regression as NN without hidden layer
 - Slope as weight
 - Intercept as bias



Activation functions

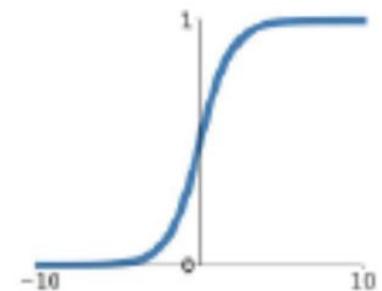
- Rectified linear unit (ReLU) is a neuron, in which the output signal appears when the inputs exceeded specified threshold

$$z = b + \sum_i x_i w_i; y = \begin{cases} z, & \text{if } z \geq 0, \\ 0, & \text{if } z < 0 \end{cases}$$



- In shallow NN sigmoid neurons are used

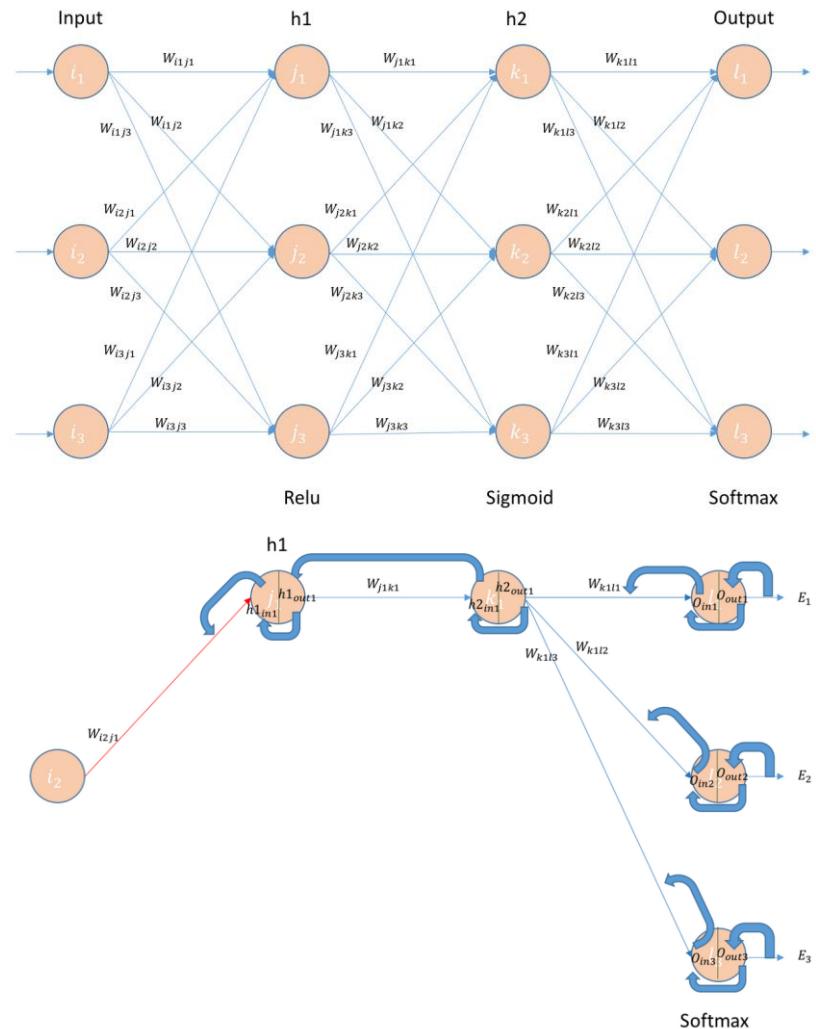
$$z = b + \sum_i x_i w_i; y = \frac{1}{1 + e^{-z}}$$



- Both allows artificial neural network to learn using backpropagation method

Backward propagation of errors

- Learning repeats two phase cycle, propagation and weight update
- When an input vector is presented to the network, it is propagated forward through the network, layer by layer, until it reaches the output layer
- The predicted output is then compared to the desired (true) one, using a loss function, and an error value is calculated
- The error values are then propagated backwards using chain rule



Home  x localhost:8888/notebooks/Image classification with Microsoft Cognitive Toolkit (autosaved) 140% ... Szukaj Logout

jupyter Image classification with Microsoft Cognitive Toolkit (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

In [4]:

```
# Train and evaluate the network.
#
def train_and_evaluate(reader_train, reader_test, max_epochs, model_func):
    # Input variables denoting the features and label data
    input_var = C.input_variable((num_channels, image_height, image_width))
    label_var = C.input_variable((num_classes))

    # Normalize the input
    feature_scale = 1.0 / 256.0
    input_var_norm = C.element_times(feature_scale, input_var)

    # apply model to input
    z = model_func(input_var_norm, out_dims=10)

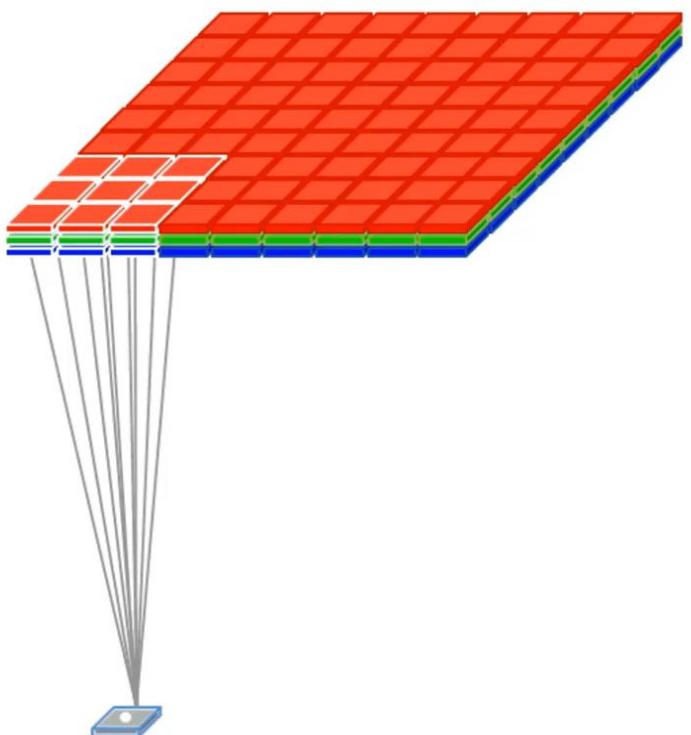
    #
    # Training action
    #

    # loss and metric
    ce = C.cross_entropy_with_softmax(z, label_var)
    pe = C.classification_error(z, label_var)

    # training config
    epoch_size      = 50000
    minibatch_size  = 32

    # Set training parameters
    lr_per_minibatch      = C.learning_rate_schedule([0.01]*6 + [0.003]*3 +
                                                       C.UnitType.minibatch,
    momentum_time_constant = C.momentum_as_time_constant_schedule(minibatch)
    l2_reg_weight          = 0.001
```

Convolutional Neural Networks



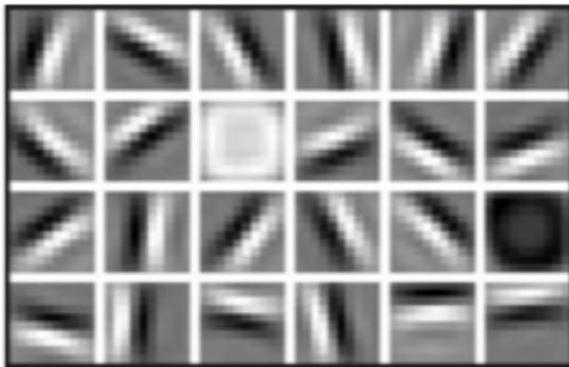
Neural Networks and images

- Regular Neural Nets don't scale well to full images
 - For 32x32x3 images a single fully-connected neuron in a first hidden layer of a regular NN would have $32*32*3 = 3072$ weights
 - 200x200x3, images would lead to neural net that have $200*200*3 = 120\ 000$ weights in a single layer
- Regular Neural Nets don't preserve image spatial structure
 - Each image is unfolded into single vector
 - For 32x32x3 images we got a vector of 3 072 digits
 - For 200x200x images we got a vector of 120 000 digits

Why Deep Learning?

- › Hand engineered features are time consuming, brittle and not scalable in practice
- › Can we learn the underlying features directly from data?

Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

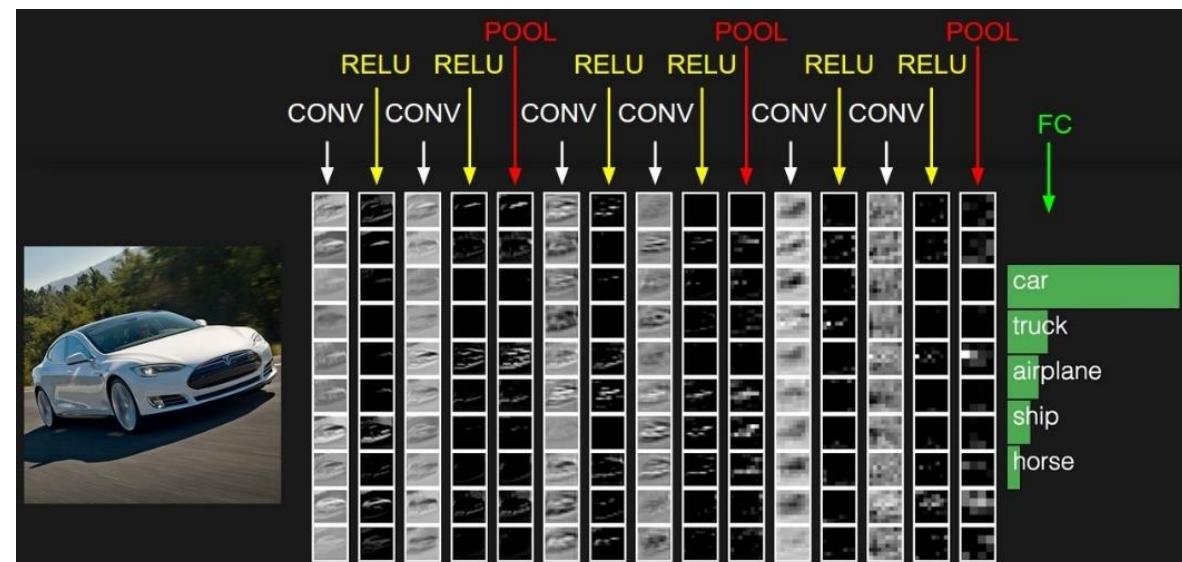
High Level Features



Facial Structure

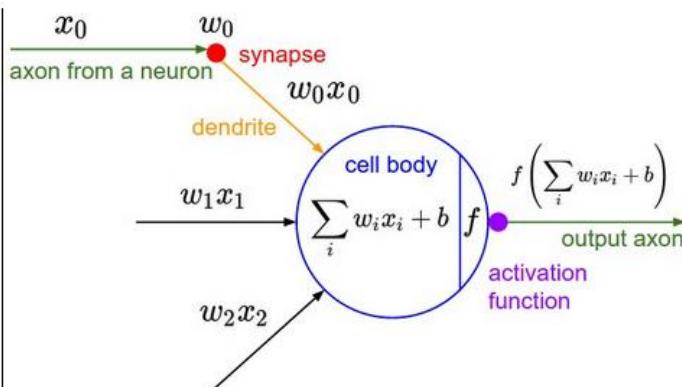
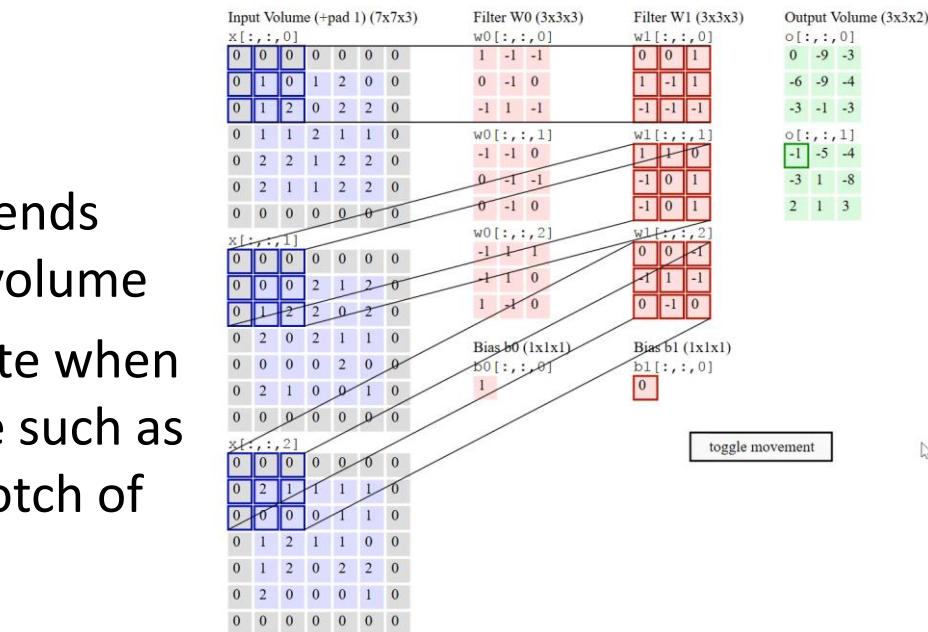
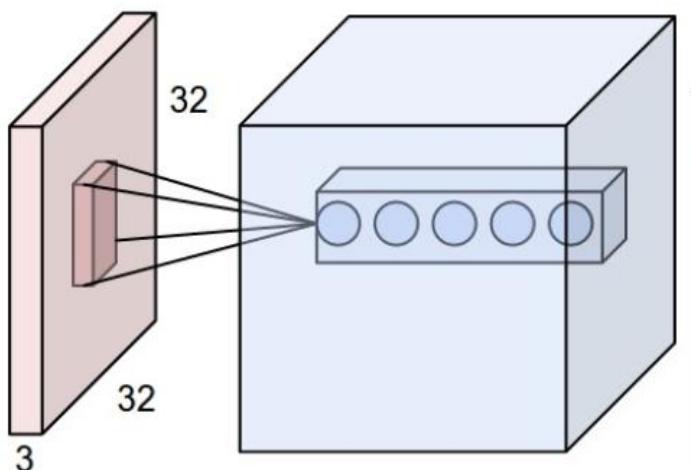
Convolutional Neural Networks

- Layers of a ConvNet have neurons arranged in 3 dimensions: width, height, depth
- Neurons in a layer are connected only to a small region of the layer before it
- Three main types of layers
 - Convolutional Layer
 - Pooling Layer
 - Fully-Connected Layer



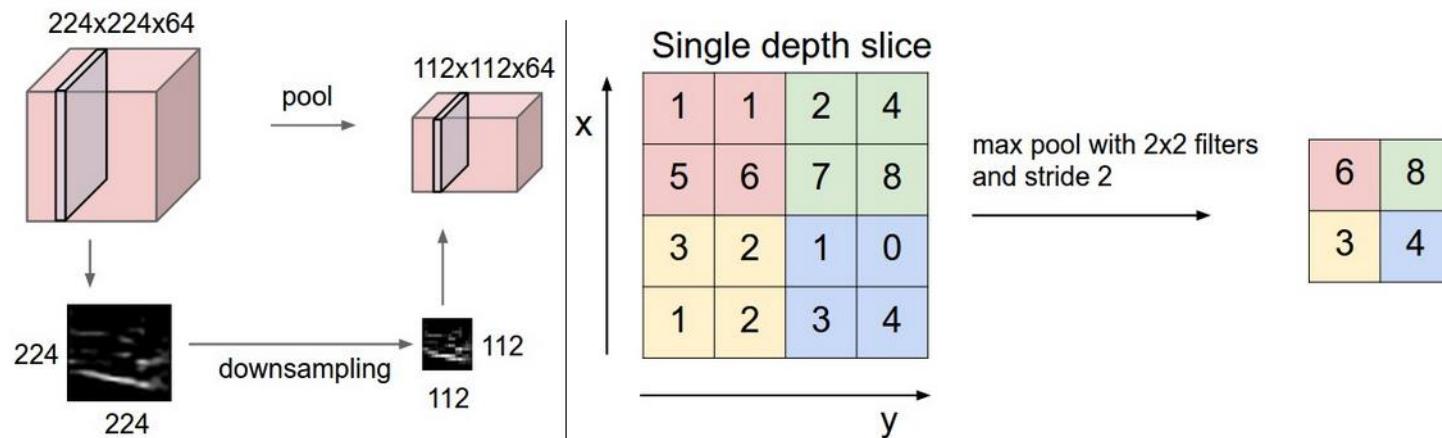
Convolutional layer

- Set of learnable filters
- Every filter is small spatially, but extends through the full depth of the input volume
- Network will learn filters that activate when they see some type of visual feature such as an edge of some orientation or a blotch of some color



Pooling layer

- Its function is to progressively reduce the spatial size of the representation to reduce the amount of parameters and computation in the network, and hence to also control overfitting
- Pooling layer operates independently on every depth slice of the input and resizes it spatially, using the MAX operation



Home Image classification with Micro +

localhost:8888/notebooks/Image classification with Micro 140% ... Szukaj Logout

jupyter Image classification with Microsoft Cognitive Toolkit (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

In [7]: `test_file = os.path.join(data_path, "gtr.png")
Image(filename=test_file)`

Out[7]:



In [8]: `from PIL import Image as PILImage
myimg = PILImage.open(test_file)
myimg = myimg.resize((32, 32), PILImage.ANTIALIAS)
myimg = np.array(myimg, dtype=np.float32)`

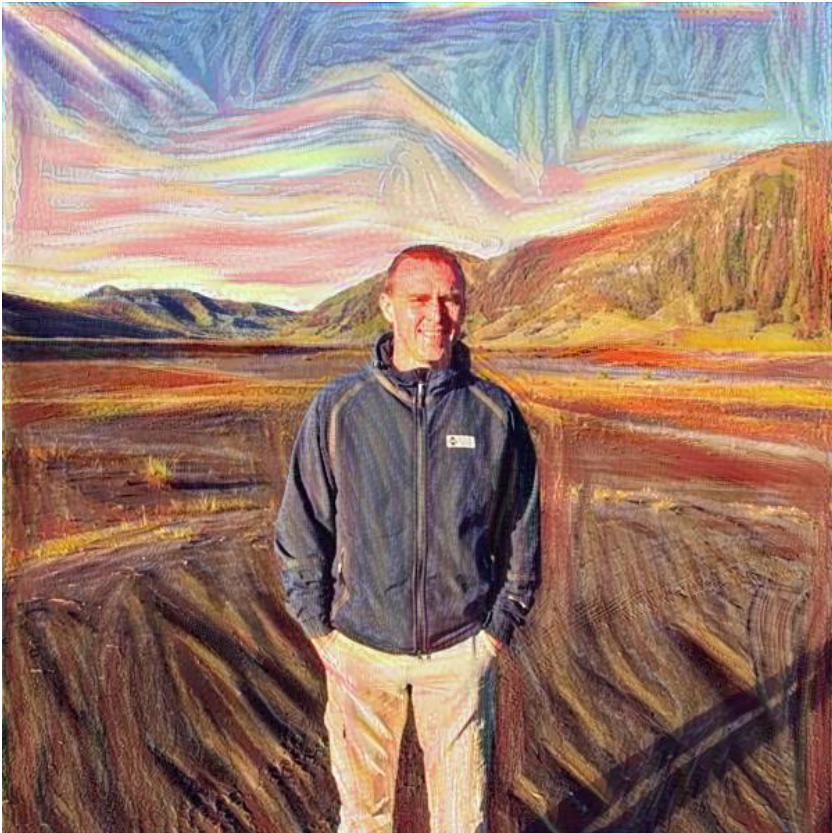
In [9]: `def eval(pred_op, image_data):
 label_lookup = ["airplane", "automobile", "bird", "cat", "deer", "dog",
 image_mean = 128.0
 image_data -= image_mean
 image_data = np.ascontiguousarray(np.transpose(image_data, (2, 0, 1)))

 result = np.squeeze(pred_op.eval({pred_op.arguments[0]: [image_data]}))

 # Return top 3 results:
 top_count = 3
 result_indices = (-np.array(result)).argsort()[:top_count]

 print("Top 3 predictions:")`

Image classification with Microsoft Cognitive Toolkit



marcin.szeliga@datacommunity.pl
<https://www.linkedin.com/in/marcinszeliga>

- Get Andrew Ng course <http://www.coursera.org/learn/neural-networks-deep-learning>
- Practice using TensorFlow <http://www.kaggle.com/learn/deep-learning>
- Read „Deep Learning” by Bengio Yoshua, Courville Aaron, Goodfellow Ian
- Watch MIT course 6.S191: Introduction to Deep Learning <http://introtodeeplearning.com/>